

2. Variational AutoEncoders

M. Ravasi

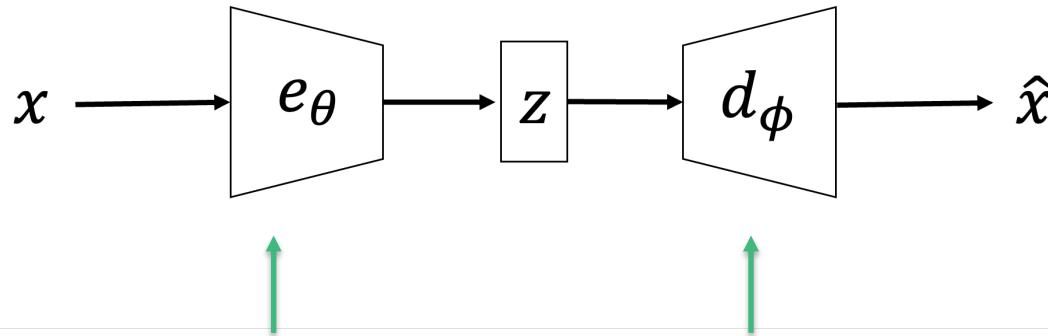
AI Summer School @ KAUST

Autoencoders

*Family of neural networks for which the **input is the same as the output**. They work by compressing the input into a latent-space representation, and then reconstructing the output from this representation.*

Autoencoders

*Family of neural networks for which the **input is the same as the output**. They work by compressing the input into a latent-space representation, and then reconstructing the output from this representation.*

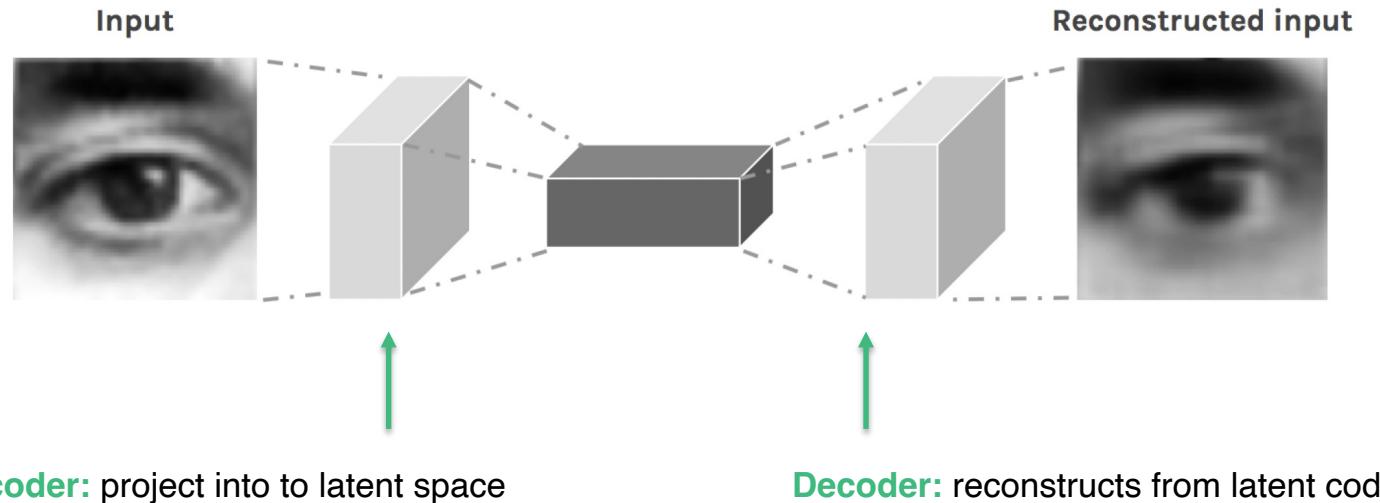


Encoder: project into latent space

Decoder: reconstructs from latent code

Autoencoders

*Family of neural networks for which the **input is the same as the output**. They work by compressing the input into a latent-space representation, and then reconstructing the output from this representation.*



Autoencoders

*Family of neural networks for which the **input is the same as the output**. They work by compressing the input into a latent-space representation, and then reconstructing the output from this representation.*

<https://douglasduhaime.com/posts/visualizing-latent-spaces.html>

Autoencoders

Mathematically described as:

$$\hat{\mathbf{x}} = d_\phi(e_\theta(\mathbf{x}))$$

The free-parameters of the 2 networks are optimized by minimizing the following loss function:

$$\hat{e}_\theta, \hat{d}_\phi = \operatorname{argmin}_{e_\theta, d_\phi} \frac{1}{N_s} \sum_i \mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)})))$$

Autoencoders - applications

Data compression

ORIGINAL
1000 x 1500, **100kb**



Instead of requesting a full-sized image, G+ requests just 1/4th the pixels...

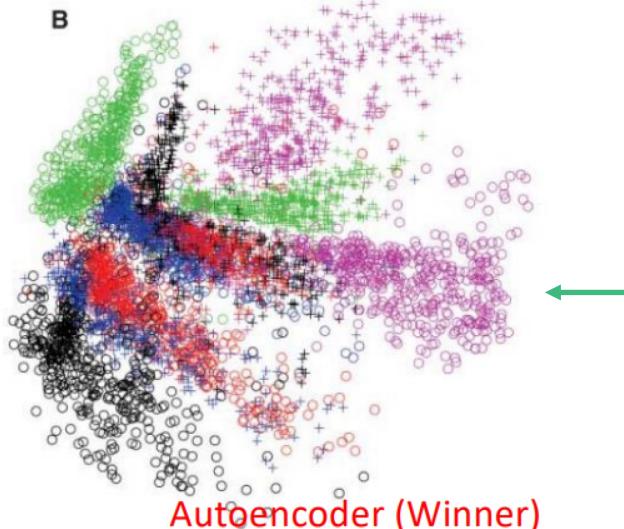
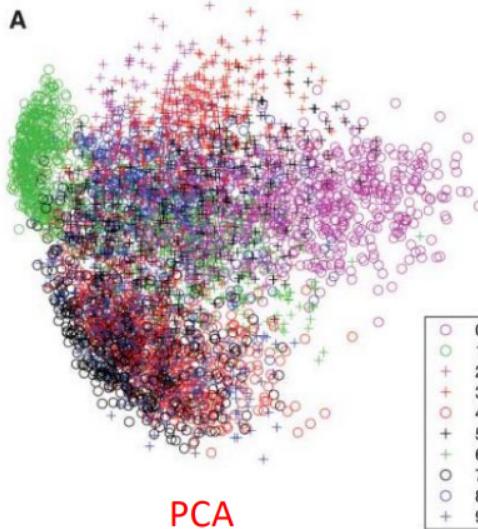
RAISR
1000 x 1500, **25kb**



...and uses RAISR to restore detail on device

Autoencoders - applications

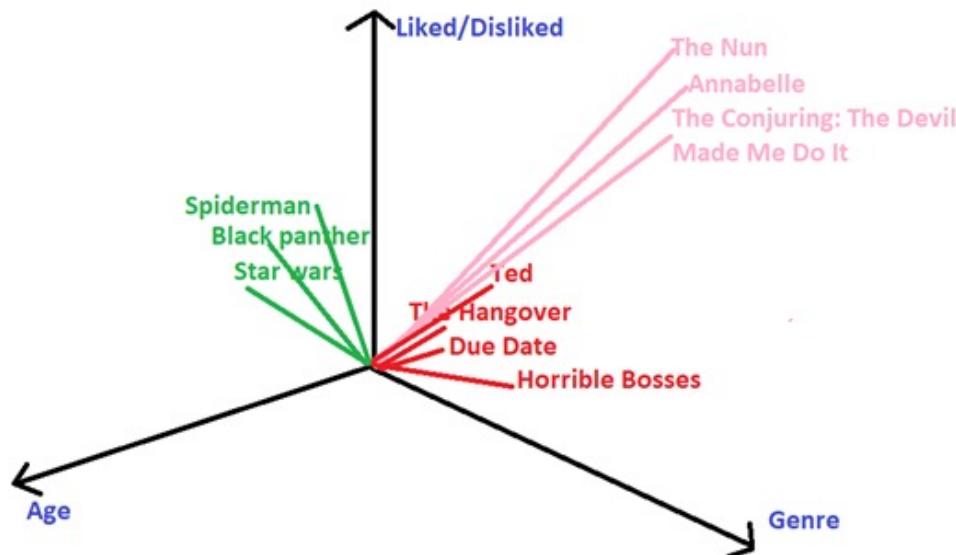
Learn robust features to be input for subsequent supervised task



Easier to do classification

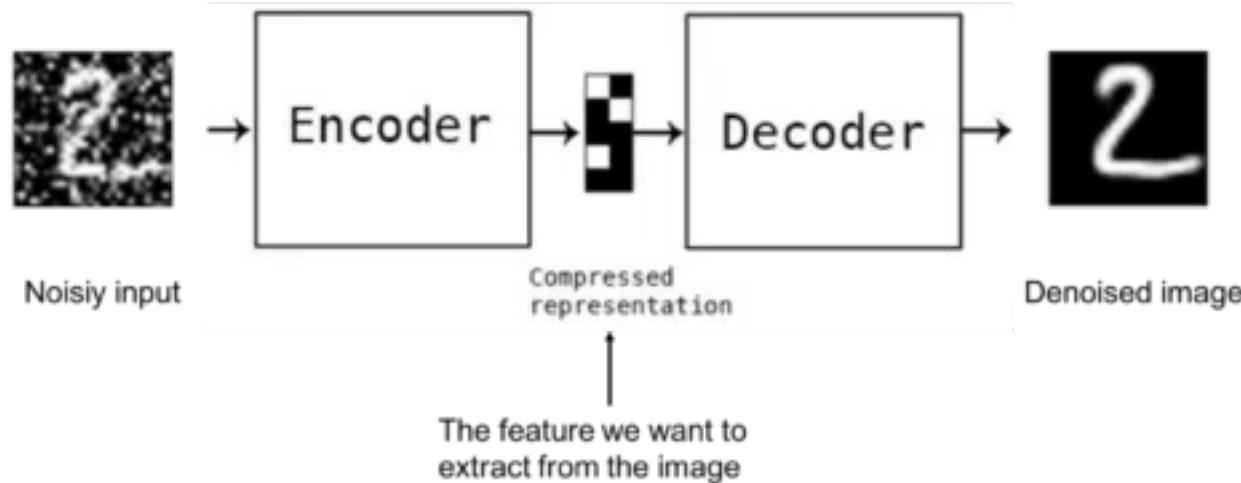
Autoencoders - applications

Perform vector math in latent space (easier to compare vectors...)

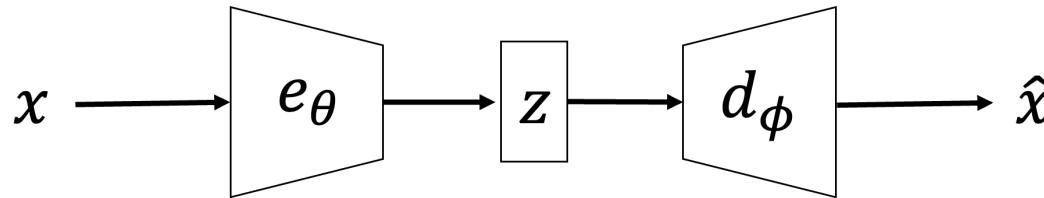


Autoencoders - applications

Denoising



Under- vs Overcomplete Autoencoders

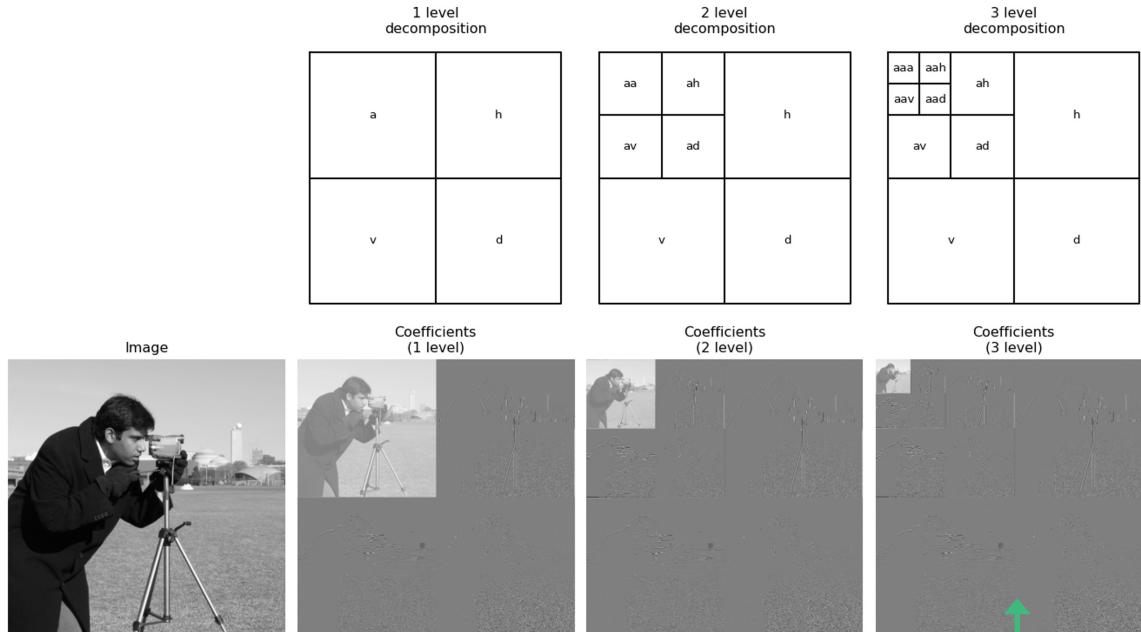


$N_z \ll N_x$ Undercomplete \rightarrow risk of high compression loss

$N_x \approx N_z$ Overcomplete \rightarrow risk of overfitting (learn identity)

→ **Need for regularization**

Overcomplete Transforms



Source: <https://medium.com/@koushikc2000/2d-discrete-wavelet-transformation-and-its-applications-in-digital-image-processing-using-matlab-1f5c68672de3>

Sparse representation

Overcomplete AEs

Add regularization to loss – **Sparse Autoencoders**

$$\mathcal{L}_r = \frac{1}{N_s} \sum_i \left(\mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)}))) + \lambda R(\mathbf{x}^{(i)}; \theta, \phi) \right)$$

Overcomplete AEs

Add regularization to loss – **Sparse Autoencoders**

$$\mathcal{L}_r = \frac{1}{N_s} \sum_i \left(\mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)}))) + \lambda R(\mathbf{x}^{(i)}; \theta, \phi) \right)$$

- L1 on latent space: $R(\mathbf{x}^{(i)}; \theta, \phi) = ||e_\theta(\mathbf{x}^{(i)})||_1$

Overcomplete AEs

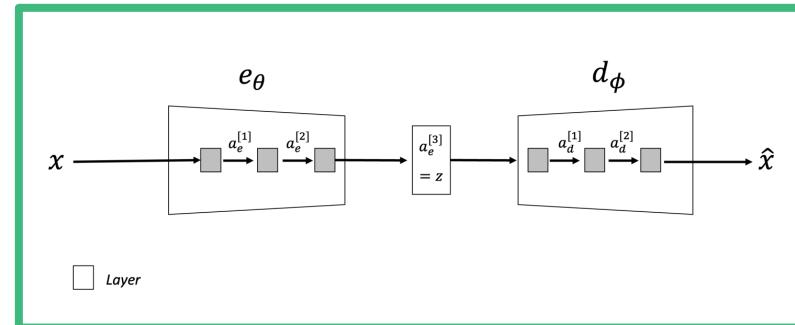
Add regularization to loss – **Sparse Autoencoders**

$$\mathcal{L}_r = \frac{1}{N_s} \sum_i \left(\mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)}))) + \lambda R(\mathbf{x}^{(i)}; \theta, \phi) \right)$$

- L1 on latent space: $R(\mathbf{x}^{(i)}; \theta, \phi) = ||e_\theta(\mathbf{x}^{(i)})||_1$

- L1 on all activations:

$$R(\mathbf{x}^{(i)}; \theta, \phi) = \sum_j ||a_e^{[j](i)}||_1 + \sum_j ||a_d^{[j](i)}||_1$$



Overcomplete AEs

Add regularization to loss – **Contractive Autoencoders**

$$\mathcal{L}_r = \frac{1}{N_s} \sum_i \left(\mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)}))) + \lambda R(\mathbf{x}^{(i)}; \theta, \phi) \right)$$

- Derivative of the latent vector over the input:

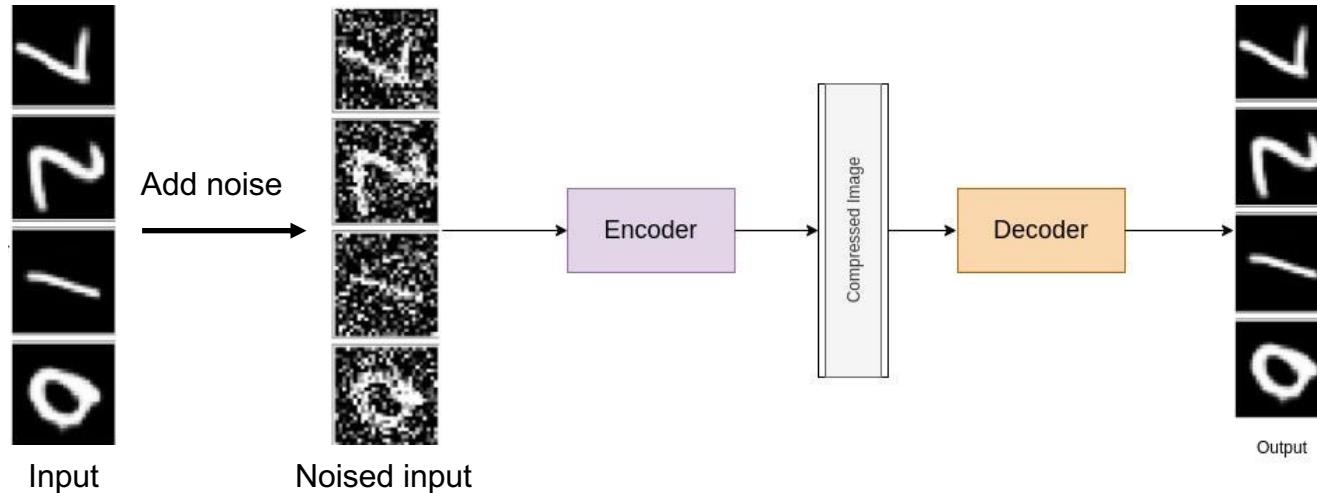
$$R(\mathbf{x}; \theta, \phi) = \|\nabla_{\mathbf{x}} \mathbf{z}\|_F$$



encourages robust latent vectors that have small sensitivity to small perturbations of the input

Overcomplete AEs

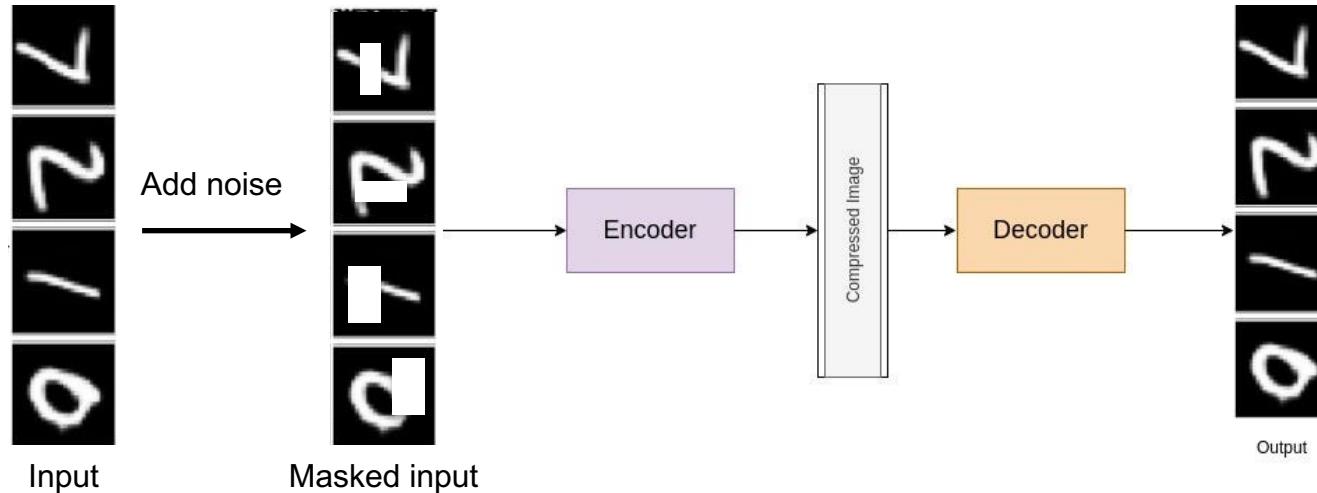
Corrupt input – Denoising Autoencoders



$$\tilde{\mathbf{x}}^{(i)} = \mathbf{x}^{(i)} + \mathbf{n}^{(i)} \quad \forall i$$

Overcomplete AEs

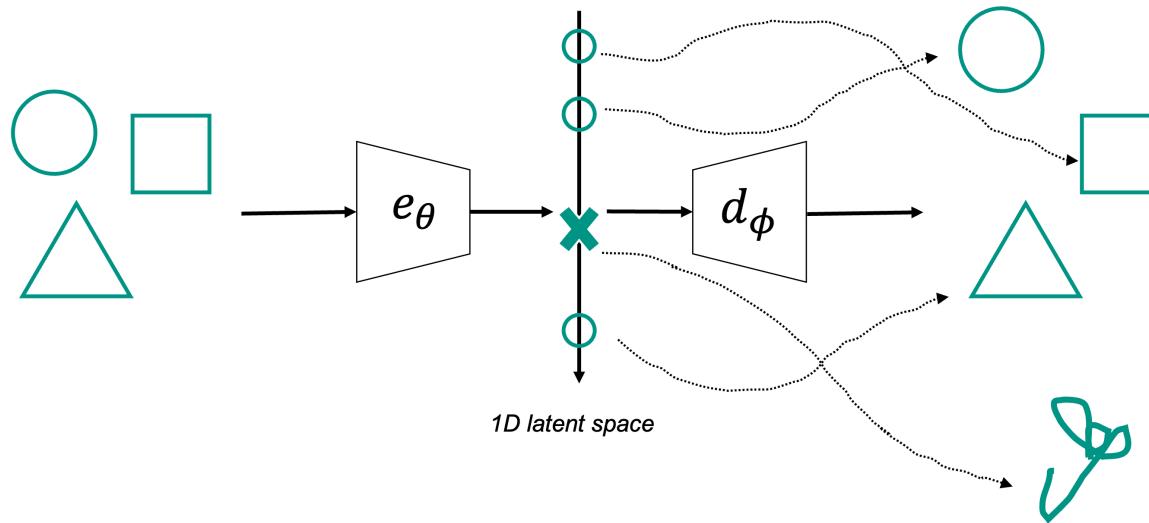
Corrupt input – Denoising Autoencoders



$$\tilde{\mathbf{x}}^{(i)} = \mathbf{x}^{(i)} \odot \mathbf{m}^{(i)} \quad \forall i$$

Autoencoders as generative models

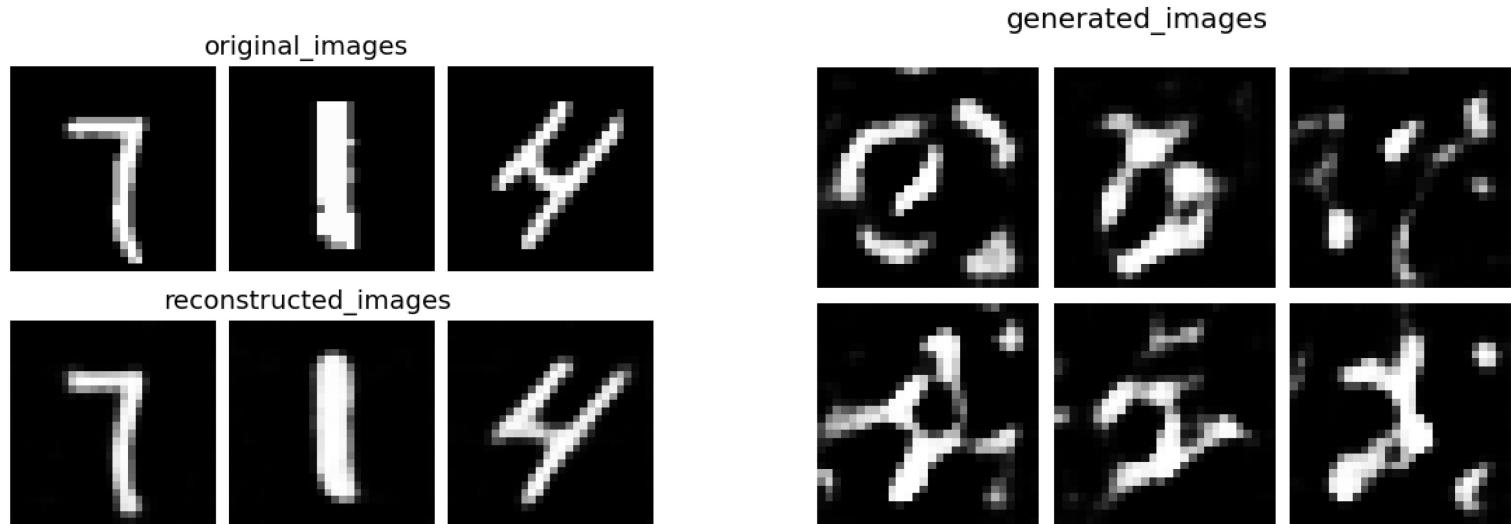
What if we **sample a new embedding vector** from z and then have the decoder reconstruct the “input” from it?



*Latent space is too **discontinuous***

Autoencoders as generative models

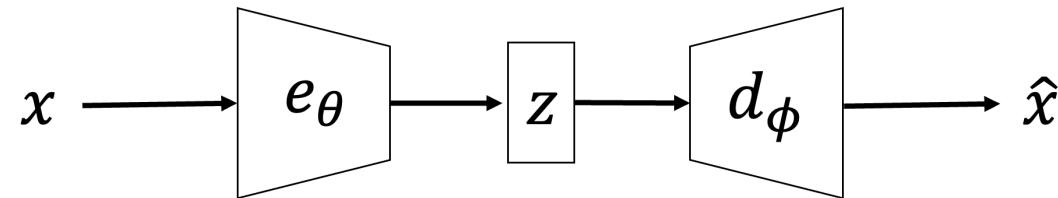
What if we **sample a new embedding vector** from z and then have the decoder reconstruct the “input” from it?



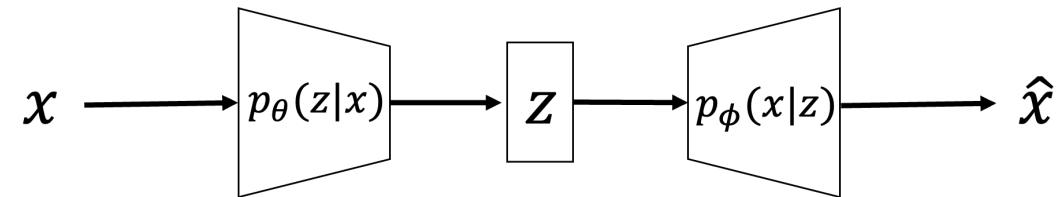
Latent space is too discontinuous

Variational Autoencoders (VAEs)

AE

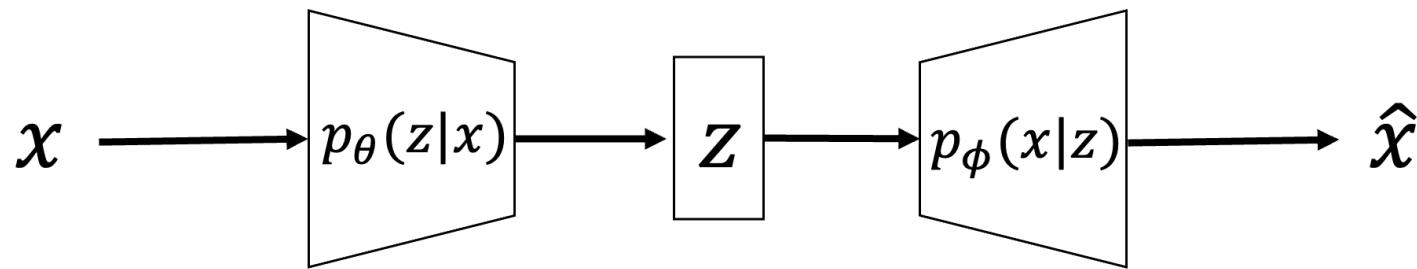


VAE



VAEs key addition: latent space sampling

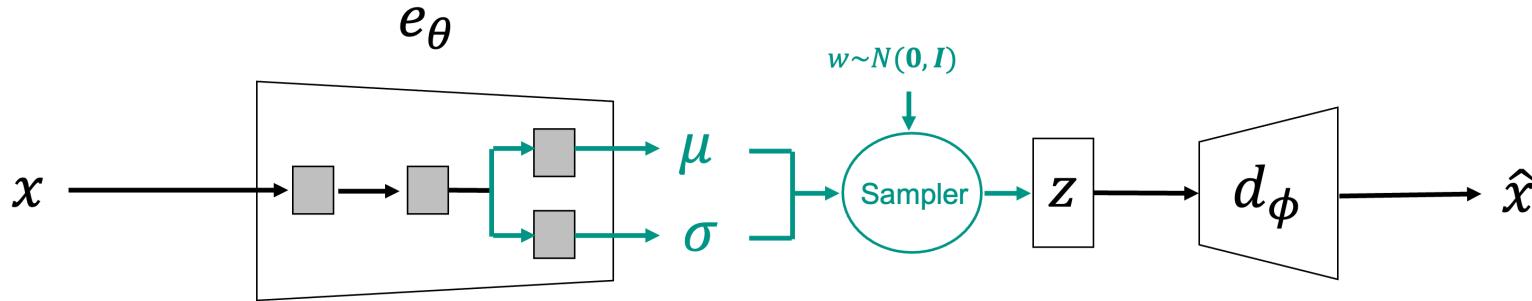
Variational AE add a probabilistic spin to AE: allow *sampling the latent space*



The two networks learn to sample some conditional probabilities...

VAEs key addition: latent space sampling

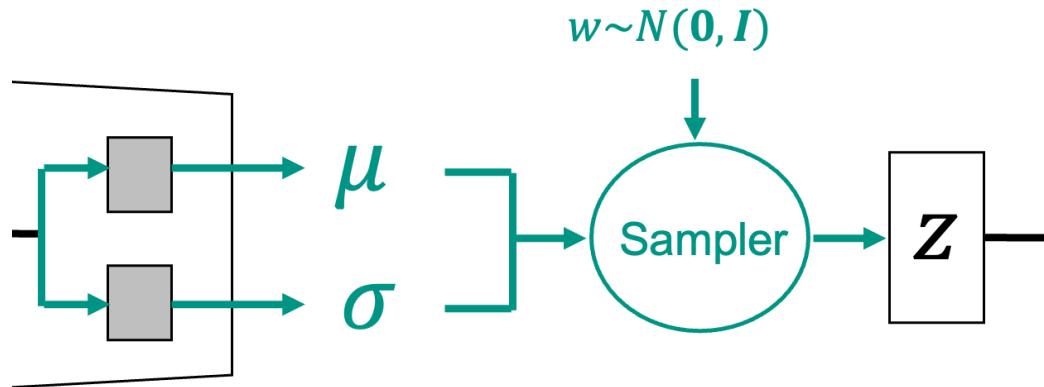
Variational AE add a probabilistic spin to AE: allow *sampling the latent space*



- 2 key changes:
- Encoding process
 - Learning process (i.e., loss)

VAEs key addition: latent space sampling

Variational AE add a probabilistic spin to AE: allow *sampling the latent space*

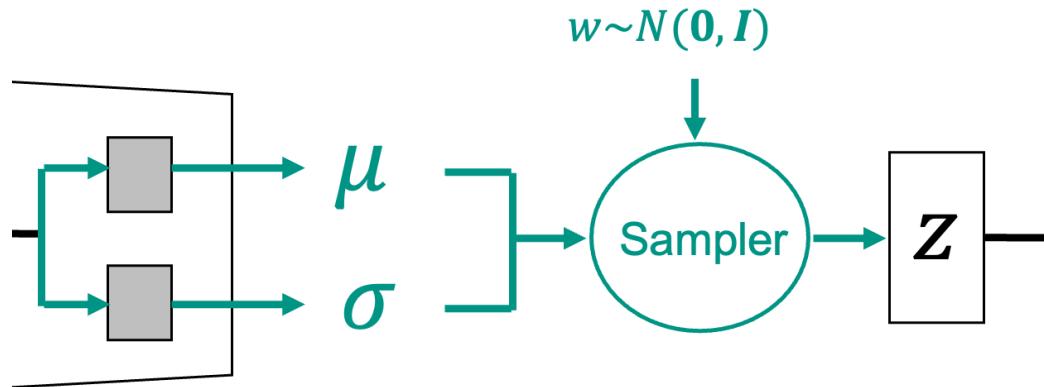


$q_{\theta}(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$: Proposal distribution

$p(\mathbf{z})$: Prior distribution

VAEs key addition: latent space sampling

Variational AE add a probabilistic spin to AE: allow *sampling the latent space*



$q_{\theta}(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$: Proposal distribution

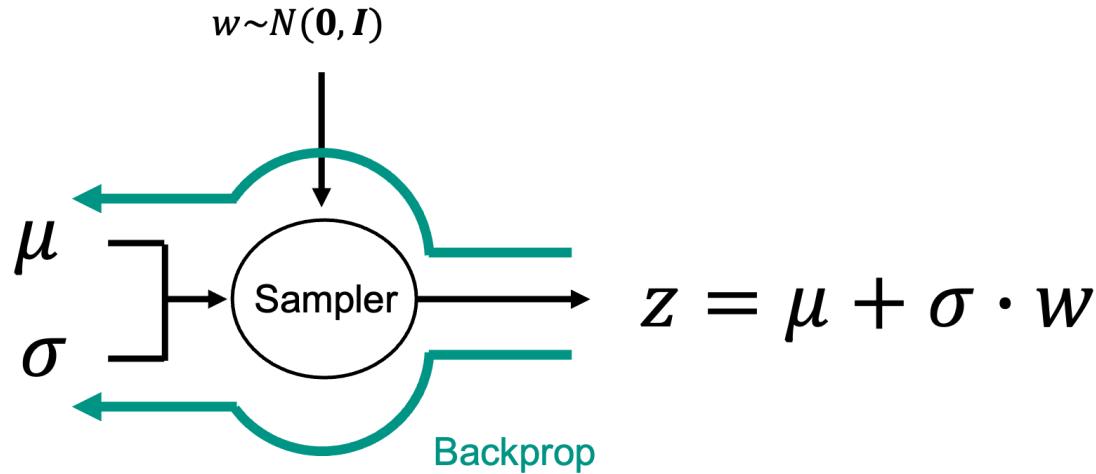
The learning process for θ changes...

$p(\mathbf{z})$: Prior distribution

Free to choose ourselves
(usually normal distribution)

Reparametrization trick

Something is required to sample a distribution, whilst still being able to do backprop



$$\frac{\partial z}{\partial \mu} = 1$$

$$\frac{\partial z}{\partial \sigma} = w$$

VAEs key addition: loss function

Add ‘probability’ regularization to loss

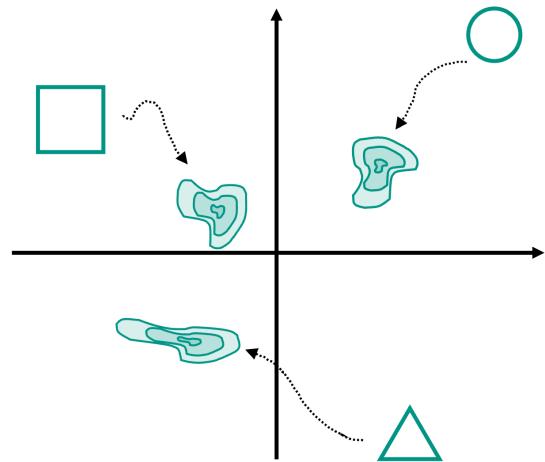
$$\mathcal{L}_r = \frac{1}{N_s} \sum_i \left(\mathcal{L}(\mathbf{x}^{(i)}, d_\phi(e_\theta(\mathbf{x}^{(i)}))) + KL(p(\mathbf{z}^{(i)}|\mathbf{x}^{(i)})||p(\mathbf{z}^{(i)})) \right)$$



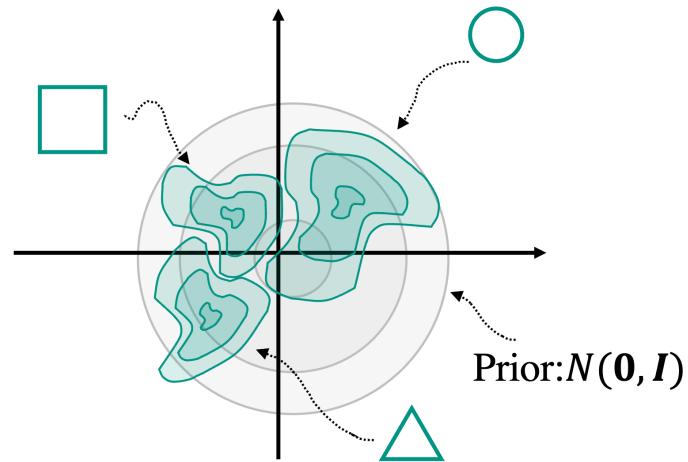
Kullback-Leibler divergence

VAEs key addition: loss function

Add ‘probability’ regularization to loss



AE latent space



VAE latent space

Kullback-Leibler divergence

Measure of closeness between two probability distributions

$$D_{KL}(Q||P) = E_{x \sim Q} \left[\log \frac{q(x)}{p(x)} \right] = \int q(x) \log \frac{q(x)}{p(x)} dx$$

- **Assymmetric metric**, not really a distance metric.
- **KL = 0**: similar, if not the same, behavior of two different distributions.
- **KL = 1**: two distributions behave in such a different manner that the expectation given the first distribution approaches zero.

Kullback-Leibler divergence

Measure of closeness between two probability distributions

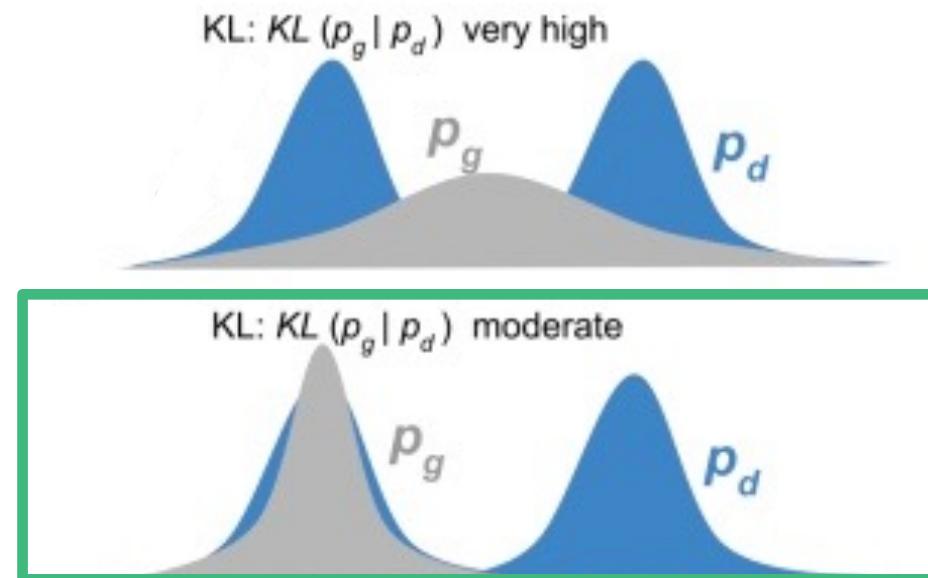
$$D_{KL}(Q||P) = E_{x \sim Q}[\log q(x)] - E_{x \sim Q}[\log p(x)]$$

- **Assymmetric metric**, not really a distance metric.
- **KL = 0**: similar, if not the same, behavior of two different distributions.
- **KL = 1**: two distributions behave in such a different manner that the expectation given the first distribution approaches zero.

Kullback-Leibler divergence

Measure of closeness between two probability distributions

p_d : target distr.
 p_g : matching distr.



This is what we will aim to achieve...

Variational Inference

Mathematical framework to **optimize parameters of one distribution to match another distribution** (usually impossible to sample and evaluate)

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z}|\mathbf{x}))$$



Bayes Rule: $p(\mathbf{z}|\mathbf{x}) = \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}$

Variational Inference (derivation)

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z}|\mathbf{x}))$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} \left[\log \left(\frac{q_{\theta}(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right) \right] =$$

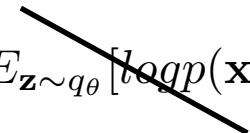
$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{z}|\mathbf{x})] =$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} \left[\log \left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})} \right) \right] =$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x}|\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{z})] + E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x})] =$$

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z})) - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x}|\mathbf{z})]$$

Evidence: No dependance on θ



Variational Inference (derivation)

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z}|\mathbf{x}))$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} \left[\log \left(\frac{q_{\theta}(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right) \right] =$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{z}|\mathbf{x})] =$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} \left[\log \left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})} \right) \right] =$$

$$\operatorname{argmin}_{\theta} E_{\mathbf{z} \sim q_{\theta}} [\log q_{\theta}(\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x}|\mathbf{z})] - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{z})] + E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x})] =$$

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z})) - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x}|\mathbf{z})]$$

Evidence: No dependance on θ



KL on prior (tractable)

Classic ML estimator on decoder (tractable)

Evidence Lower BOund (ELBO)

A different interpretation of learning process

$$\log p(\mathbf{x}) - KL(q_\theta(\mathbf{z}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q_\theta} [\log p(\mathbf{x}|\mathbf{z})] - KL(q_\theta(\mathbf{z}) || p(\mathbf{z}))$$

Probability we want to maximize in
generative modelling

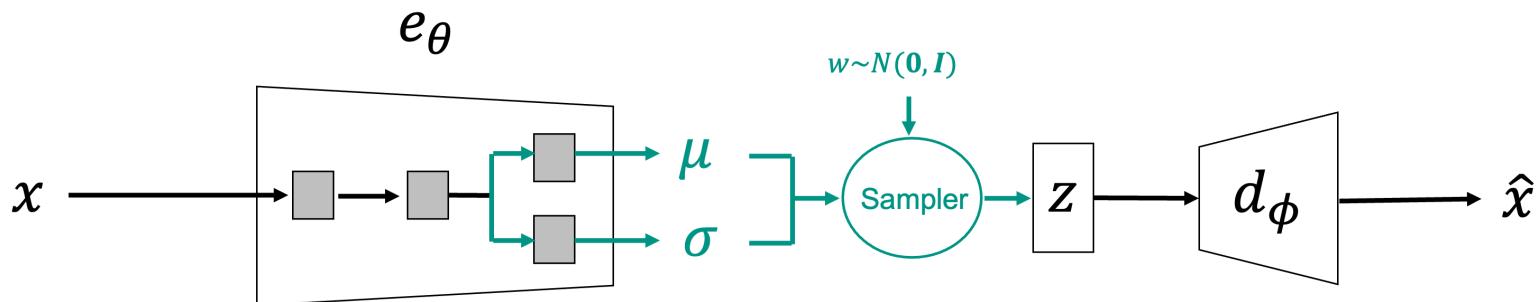
What we actually maximise

Since KL is always positive (lower bound on evidence):

$$\log p(\mathbf{x}) - KL(q_\theta(\mathbf{z}) || p(\mathbf{z}|\mathbf{x})) \leq \log p(\mathbf{x})$$

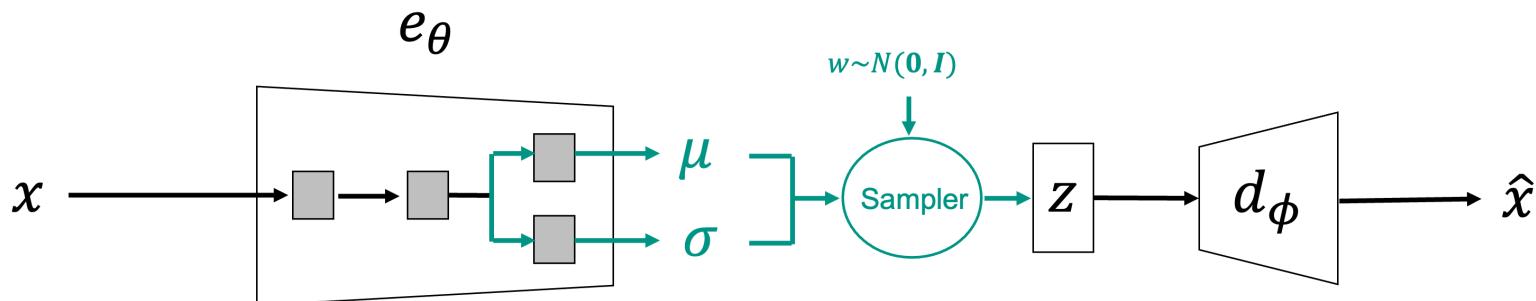
VAE in practice

$$\underset{\theta}{\operatorname{argmin}} \text{ } KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z})) - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x} | \mathbf{z})]$$
$$q_{\theta}(\mathbf{z}) \sim \mathcal{N}(e_{\theta, \mu}(\mathbf{x}), \text{diag}(e_{\theta, \sigma}(\mathbf{x}))) \quad p(\mathbf{z}) \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$



VAE in practice

$$\operatorname{argmin}_{\theta} KL(q_{\theta}(\mathbf{z}) || p(\mathbf{z})) - E_{\mathbf{z} \sim q_{\theta}} [\log p(\mathbf{x} | \mathbf{z})]$$



VAE in practice

Learning Algorithm

1. Select a batch of training samples and feed to encoder
2. Divide the output of the encoder into two parts: mean and std (each of them has the same number of samples as the input batch)
3. Sample a latent space realization per training sample
4. Feed each latent vector into the decoder
5. Compute the loss: 1) MSE of reconstruction, 2) KL between prior and proposal using the mean and std obtained in step2. + backprop.

VAE in practice

Sampling Algorithm

1. Sample a latent space realization
2. Feed each latent vector into the decoder

Encoder is not used in inference!